

Chinese CEFR Labels Inherit an Ambiguity: 46% of Texts Change Band with the Source Standard

Alvaro Serrano

Independent researcher

alvaro.serrano.gp@gmail.com

orcid.org/0009-0006-4701-9026

Abstract

Multilingual CEFR resources now span thirteen languages (Imperial et al., 2025); Chinese is not among them. The obvious route in is HSK, the Chinese proficiency scale, which is routinely mapped onto the CEFR’s six bands. That route is less well defined than it looks. “HSK 3.0” names two official documents — the GF0025-2021 national grading standard and the examination syllabus in force since July 2026 — which assign different levels to 41.5% of the vocabulary they share (Serrano, 2026). We show the ambiguity propagates: grading 102 graded readers against one document rather than the other moves **46.1%** of them into a different CEFR band under the usual one-to-one correspondence, and **44.1%** under a compressed correspondence that assumes $\text{HSK } 6 \approx \text{B2}$. It replicates on a disjoint 30-text held-out split (46.7%, 43.3%). The movement is not noise: **42 of 47** texts move *upward*, and on the held-out split all 14 do. Choosing the wrong document does not add error, it shifts the distribution. We then test the mapping itself against a reference the field has not used. Taiwan’s TOCFL carries an *official* CEFR correspondence, so labelling the same texts through it rather than through HSK gives a check on the folk tables. They disagree on **95.1%** of texts, and the direction is absolute: the HSK route places every disagreeing text higher, typically by one to two bands. Adopting the compressed correspondence halves the disagreement, to 50.0% — the first quantitative support we know of for the practitioner claim that HSK levels overstate CEFR. We argue that any Chinese entry in a CEFR corpus must record which Chinese standard its labels derive from, and that the standard should be the one that has an alignment.

1 The gap, and why it is awkward to fill

Imperial et al. (2025) assemble 494,491 CEFR-labelled texts across thirteen languages and standardise them into one schema. Chinese is absent. This is not an oversight of scale — Chinese is among the most-studied second languages in the world — but a consequence of the inclusion criteria, which require a permissive licence, CEFR labels produced or validated by domain experts, and human authorship.

The natural way to produce Chinese CEFR labels at scale is indirect: grade text against HSK, then map HSK levels onto CEFR bands. Learner-facing material does this constantly. The mapping is not official — China’s own standard declines to align itself to the CEFR — but it is ubiquitous, and it is the implicit basis on which Chinese material is described as “B1” or “C1” in practice.

This note reports an obstacle that sits before the mapping, and that a better mapping cannot fix.

2 Two documents named HSK 3.0

The name “HSK 3.0” is applied to two distinct official publications:

- 《国际中文教育中文水平等级标准》 (GF0025-2021), the national grading standard, in force since 1 July 2021. Nine levels, 11,092 words, 3,000 characters.
- 新版 HSK 考试大纲, the examination syllabus, published November 2025 and in force since July 2026. 10,896 words, 3,088 recognition characters, and a separate 1,200-character writing inventory.

Serrano (2026) compares them directly: of the 9,674 words the two share, 41.5% sit at

different levels; of the 2,945 shared characters, 40.7% do. Grading a corpus of authentic graded readers against one rather than the other changes the assigned HSK level of 48.0% of texts. Tools and reference sites that report “HSK 3.0” levels overwhelmingly do not say which document they used.

The question this note asks is narrow: does that disagreement survive projection onto the CEFR’s coarser six-band scale, or does coarsening wash it out?

3 Method

Corpus. 102 graded readers spanning six difficulty shelves, with per-word segmentation authored rather than produced by a segmenter, plus a disjoint 30-text held-out split. Both are released CC BY 4.0 with a datasheet. We grade the authored segmentation directly; re-tokenising the raw strings with a segmenter yields materially different numbers and is not comparable to the published figures.

Grading. Each text is graded twice with the same open implementation (Serrano, 2026), once against GF0025-2021 and once against the 2025 syllabus. A text’s level is the lowest level whose cumulative character inventory covers 95% of its running characters, the minimal coverage threshold from Laufer and Ravenhorst-Kalovski (2010). Nothing about the procedure differs between the two runs except the inventory.

Correspondences. We use two, and report every figure under both. The **naïve** correspondence is the one-to-one reading ubiquitous in learner material: HSK n maps to the n th CEFR band, with HSK 3.0’s combined levels 7–9 mapped to C2. The **compressed** correspondence encodes the objection practitioners raise — that vocabulary coverage overstates communicative range, so HSK 6 lands nearer B2 than C2 — by mapping $\{1, 2\} \rightarrow A1$, $3 \rightarrow A2$, $4 \rightarrow B1$, $\{5, 6\} \rightarrow B2$, $\{7, 8\} \rightarrow C1$, $9 \rightarrow C2$.

Neither is authoritative. That is the point of using both: a finding that holds under both does not rest on the choice.

4 Results

Table 1 gives the result. Coarsening to six bands does not wash the disagreement out:

	Readers ($n=102$)	Held-out ($n=30$)
HSK level changes	49 (48.0%)	14 (46.7%)
CEFR band, naïve	47 (46.1%)	14 (46.7%)
upward	42	14
downward	5	0
CEFR band, compressed	45 (44.1%)	13 (43.3%)
upward	42	13
downward	3	0

Table 1: Texts whose band differs depending on which HSK document was used. “Upward” means the 2025 syllabus placed the text higher than the 2021 standard did.

	A1	A2	B1	B2	C1	C2
2021 basis	3	28	32	18	14	7
2025 basis	5	10	29	32	19	7

Table 2: Band distribution of the same 102 texts under the naïve correspondence. The A2 population falls by 64% and the B2 population rises by 78%; the corpus is not relabelled at random, it is relabelled harder.

46.1% of texts change band under the naïve correspondence, against 48.0% that change HSK level. Compressing the correspondence — which merges adjacent HSK levels and so should absorb single-level shifts — reduces it only to 44.1%. The pattern replicates on the held-out split at 46.7% and 43.3%.

The direction is the more consequential finding. Of the 47 texts that move under the naïve correspondence, 42 move upward; on the held-out split, all 14 do. Table 2 shows the effect on the corpus as a whole: the A2 population falls from 28 to 10 while B2 rises from 18 to 32. A researcher who picks the wrong document does not acquire noisy labels. They acquire a corpus that is systematically mis-levelled in one direction, which is the failure mode least likely to be caught by inspection and most likely to bias a model trained on it.

5 A reference the field has not used

Everything above compares two PRC documents to each other. It cannot say whether *either* produces a defensible CEFR label, because the PRC standard publishes no CEFR alignment: the correspondences in §3 are folk tables, and there is nothing to check them against.

	Readers (<i>n</i> =102)	Held-out (<i>n</i> =30)
<i>Taiwan vs. PRC, naive correspondence</i>		
label differs	97 (95.1%)	25 (83.3%)
PRC places higher	97	25
Taiwan places higher	0	0
<i>Taiwan vs. PRC, compressed correspondence</i>		
label differs	51 (50.0%)	20 (66.7%)
PRC places higher	45	18
Taiwan places higher	6	2

Table 3: Labelling the same texts through Taiwan’s officially CEFR-aligned framework rather than through HSK. Under the naive table the disagreement is total and entirely one-directional.

There is something to check them against. Chinese has a second proficiency ecosystem. Taiwan’s TOCFL was built at National Taiwan Normal University with the CEFR as its blueprint and carries an official level correspondence, and the national standard behind it, TBCL, was developed at the National Academy for Educational Research and designed to align to both CEFR and ACTFL. Its 7,517 word list is published free by SC-TOPI, banded by level. Where the PRC pair publishes neither a CEFR alignment nor a standard-to-exam correspondence, the Taiwan pair publishes both.

Method. We grade the same corpus a third time, against the TOCFL inventory, converted to simplified characters with OpenCC. Both frameworks are applied character-level at the same 95% threshold with the same proper-noun handling, so only the framework varies. TOCFL publishes no character list at this tier, so its inventory is derived from its words — a character takes the lowest level of any word containing it. To check what that derivation costs we apply it to the PRC’s own words, where an official list exists: it reproduces the official grade on **102 of 102** texts. The gap below is therefore a property of the frameworks, not of the derivation.

Result. Under the naive one-to-one table, 95.1% of texts receive a different CEFR label depending on which country’s framework produced it, and *every one* of the 97 disagreements runs the same way: the HSK route places the text higher, by one band in 52 cases, two in 38 and three in 7.

The compressed correspondence — the one practitioners argue for, in which HSK 6 lands nearer B2 than C2 — cuts the disagreement to 50.0%, raising exact agreement from 5 texts to 51. That is, to our knowledge, the first quantitative test of the long-standing claim that HSK levels overstate CEFR, and it supports it: moving the table downward moves it toward the framework that actually has an alignment. It does not vindicate it either, since half the corpus still disagrees and the residual is still 45 to 6 in the same direction.

6 What this implies for multilingual CEFR resources

The provenance field is not optional. The Imperial et al. (2025) schema records source dataset, language, granularity, production category and licence. For a Chinese entry this is insufficient: two entries could share every one of those values, be produced by the same procedure over the same texts, and still disagree on the band of nearly half of them. Chinese needs the source standard, with its version, recorded as data. We would expect the same to apply wherever a language’s CEFR labels are derived through a national framework rather than assigned directly.

Prefer the framework that has an alignment. The cleanest reading of §5 is that a Chinese CEFR label should be derived through TOCFL or TBCL, where the correspondence is published by the body that defines the levels, rather than through HSK, where it is reconstructed. This also gives a labelling route that does not depend on resolving the PRC ambiguity at all.

No correspondence table fixes the PRC ambiguity. It is tempting to read §4 as an argument for a better HSK–CEFR mapping. It is not. Both correspondences we test are applied identically to both gradings; the disagreement is generated before the mapping is consulted and passes through it. A third correspondence would behave the same way.

Derived labels are not expert labels. We are explicit that the labels here are algorithmic projections of a coverage threshold, not expert-assigned CEFR levels, and so do

not meet the gold-standard criterion of Imperial et al. (2025). The contribution is not a proposed dataset entry. It is a constraint that a future entry will have to satisfy.

7 Limitations

The corpus is LLM-assisted, disclosed in its datasheet, and pedagogically constrained; it is not a sample of naturally occurring Chinese and would not meet the human-authorship criterion of Imperial et al. (2025). Its levels are ours, not experts'. Coverage is necessary but not sufficient for comprehension: grammar, register and world knowledge are unmodelled, and the 95% threshold is imported from word-level English research rather than established for L2 Chinese characters. Neither HSK correspondence we use is official, and we make no claim that either is correct — only that the finding does not depend on which is chosen. The TOCFL comparison carries its own limits: the inventory is derived rather than official, TBCL's 3,100-character list would be the better source and we did not have it; the word list stops at 流利級 and so describes no C2 tier; and 338 of 7,517 entries collapse together under conversion to simplified characters. Finally, the 2025 syllabus takes effect in July 2026, so its practical consequences are not yet observable in learner outcomes.

8 Availability

The grading implementation is MIT licensed and installable as `pip install hsk30`; the corpus and held-out split are CC BY 4.0. Every figure in this note is regenerated by `scripts/cefr_mapping.py` and `scripts/tocfl_compare.py` in the same repository. The TOCFL list is not redistributed — SC-TOP asserts rights over it — so the second script reads a copy downloaded free from the official site, and the repository ships only the code.

Ethics and data statement

The corpus contains no personal data and describes invented characters in everyday situations; its LLM-assisted authorship is disclosed in its datasheet and again in the limitations above. Word and character lists derive from

published national standards. No human subjects were involved. An intended application is a Chinese-learning website operated by the author.

References

- Joseph Marvin Imperial, Abdullah Barayan, Regina Stodden, Rodrigo Wilkens, Ricardo Muñoz Sánchez, Lingyun Gao, Melissa Torgbi, Dawn Knight, Gail Forey, Reka R. Jablonkai, Ekaterina Kochmar, Robert Joshua Reynolds, Eugénio Ribeiro, Horacio Saggion, Elena Voldina, Sowmya Vajjala, Thomas François, Fernando Alva-Manchego, and Harish Tayyar Madabushi. 2025. [UniversalCEFR: Enabling open multilingual research on language proficiency assessment](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 9703–9755, Suzhou, China. Association for Computational Linguistics.
- Batia Laufer and Geke C. Ravenhorst-Kalovski. 2010. Lexical threshold revisited: Lexical text coverage, learners' vocabulary size and reading comprehension. *Reading in a Foreign Language*, 22(1):15–30.
- Alvaro Serrano. 2026. [Which HSK 3.0? two official documents, and half of all grading decisions change](#). Preprint.